

Yolo를 활용한 객체 인식

학습 내용

- yolo3를 설치 후, 이를 활용한 객체 인식을 수행해 봅니다.

목차

01 사전 준비 - 설치, 환경 설정

02 yolo를 활용 - 예측수행

03 yolo 예측 결과 확인

01 사전 준비 - 설치, 환경 설정

목차로 이동하기

- OpenCV: 객체 인식 및 이미지 처리를 위한 라이브러리로 설치가 필요

```
In [4]: # OpenCV 설치
!pip install opencv-python

Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/public/simple/
Requirement already satisfied: opencv-python in /usr/local/lib/python3.10/dist-packages (4.7.0.72)
Requirement already satisfied: numpy>=1.21.2 in /usr/local/lib/python3.10/dist-packages (from opencv-python) (1.22.4)
```

- Darknet: YOLO 프레임워크로서, 객체 인식을 수행하기 위해 설치가 필요

```
In [6]: # !git clone https://github.com/AlexeyAB/darknet.git
# !git clone https://github.com/pjreddie/darknet
```

```
In [10]: # darknet 파일 확인
!cd darknet/; ls; make
```

3rdparty	data	README.md
build	Deer_un.jpeg	results
build.ps1	image_yolov3.sh	scripts
cfg	image_yolov4.sh	src
cmake	include	vcpkg.json
CMakeLists.txt	json_mjpeg_streams.sh	vcpkg.json.opencv23
DarknetConfig.cmake.in	LICENSE	video_yolov3.sh
darknet_images.py	Makefile	video_yolov4.sh
darknet.py	net_cam_v3.sh	
darknet_video.py	net_cam_v4.sh	

```

mkdir -p ./obj/
mkdir -p backup
chmod +x *.sh
g++ -std=c++11 -std=c++11 -Iinclude/ -I3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-unknown-pragmas -fPIC -rdynamic -Ofast -c ./src/image_opencv.cpp -o obj/image_opencv.o
g++ -std=c++11 -std=c++11 -Iinclude/ -I3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-unknown-pragmas -fPIC -rdynamic -Ofast -c ./src/http_stream.cpp -o obj/http_stream.o
./src/http_stream.cpp: In member function 'bool JSON_sender::write(const char*)':
./src/http_stream.cpp:253:21: warning: unused variable 'n' [-Wunused-variable]
253 |             int n = _write(client, outputbuf, outlen);
    |             ^
./src/http_stream.cpp: In function 'void set_track_id(detection*, int, float, float, float, int, int, int)':
./src/http_stream.cpp:867:27: warning: comparison of integer expressions of different signedness: 'int' and 'std::vector<detection_t>::size_type' {aka 'long unsigned int'} [-Wsign-compare]
867 |         for (int i = 0; i < v.size(); ++i) {
    |                        ~~~^~~~~~
./src/http_stream.cpp:875:33: warning: comparison of integer expressions of different signedness: 'int' and 'std::vector<detection_t>::size_type' {aka 'long unsigned int'} [-Wsign-compare]
875 |         for (int old_id = 0; old_id < old_dets.size(); ++old_id) {
    |                                ~~~~~~^~~~~~
./src/http_stream.cpp:894:31: warning: comparison of integer expressions of different signedness: 'int' and 'std::vector<detection_t>::size_type' {aka 'long unsigned int'} [-Wsign-compare]
894 |         for (int index = 0; index < new_dets_num*old_dets.size(); ++index)
    |                                ~~~~~~^~~~~~
./src/http_stream.cpp:930:28: warning: comparison of integer expressions of different signedness: 'std::deque<std::vector<detection_t> >::size_type' {aka 'long unsigned int'} and 'int' [-Wsign-compare]
930 |         if (old_dets_dq.size() > deque_size) old_dets_dq.pop_front();
    |            ~~~~~~^~~~~~
gcc -Iinclude/ -I3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-unknown-pragmas -fPIC -rdynamic -Ofast -c ./src/gemm.c -o obj/gemm.o
./src/gemm.c: In function 'convolution_2d':
./src/gemm.c:2044:15: warning: unused variable 'out_w' [-Wunused-variable]
2044 |     const int out_w = (w + 2 * pad - ksize) / stride + 1;    // output_width=input_width for stride=1 and pad=1
    |               ^~~~~
./src/gemm.c:2043:15: warning: unused variable 'out_h' [-Wunused-variable]
2043 |     const int out_h = (h + 2 * pad - ksize) / stride + 1;    // output_height=input_height for stride=1 and pad=1
    |               ^~~~~
gcc -Iinclude/ -I3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-unknown-pragmas -fPIC -rdynamic -Ofast -c ./src/utils.c -o obj/utils.o
./src/utils.c: In function 'custom_hash':
./src/utils.c:1082:12: warning: suggest parentheses around assignment used as truth value [-Wparentheses]
1082 |     while(c = *str++)

```

```

|           ^
In file included from /usr/include/string.h:495,
                 from include/darknet.h:14,
                 from ./src/utils.h:3,
                 from ./src/utils.c:4:
In function 'strncpy',
    inlined from 'copy_string' at ./src/utils.c:552:5:
/usr/include/x86_64-linux-gnu/bits/string_fortified.h:106:10: warning: '_
__builtin_strncpy' specified bound depends on the length of the source argument [-Wstringop-overflow=]
   106 |     return __builtin___strncpy_chk (__dest, __src, __len, __bos (__de
st));
       |           ^~~~~~
~~~~~
./src/utils.c: In function 'copy_string':
./src/utils.c:552:22: note: length computed here
   552 |     strncpy(copy, s, strlen(s)+1);
       |           ^~~~~~
gcc -linclude/ -I3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-u
nknown-pragmas -fPIC -rdynamic -Ofast -c ./src/dark_cuda.c -o obj/dark_cuda.o
gcc -linclude/ -I3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-u
nknown-pragmas -fPIC -rdynamic -Ofast -c ./src/convolutional_layer.c -o obj/convolut
ional_layer.o
./src/convolutional_layer.c: In function 'resize_convolutional_layer':
./src/convolutional_layer.c:898:9: warning: unused variable 'old_h' [-Wunus
sed-variable]
   898 |     int old_h = l->h;
       |         ^~~~~
./src/convolutional_layer.c:897:9: warning: unused variable 'old_w' [-Wunus
sed-variable]
   897 |     int old_w = l->w;
       |         ^~~~~
./src/convolutional_layer.c: In function 'forward_convolutional_layer':
./src/convolutional_layer.c:1342:32: warning: unused variable 't_intput_siz
e' [-Wunused-variable]
   1342 |         size_t t_intput_size = binary_transpose_align_inpu
t(k, n, state.workspace, &l.t_bit_input, ldb_align, l.bit_align);
       |                                ^~~~~~
gcc -linclude/ -I3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-u
nknown-pragmas -fPIC -rdynamic -Ofast -c ./src/list.c -o obj/list.o
gcc -linclude/ -I3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-u
nknown-pragmas -fPIC -rdynamic -Ofast -c ./src/image.c -o obj/image.o
gcc -linclude/ -I3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-u
nknown-pragmas -fPIC -rdynamic -Ofast -c ./src/activations.c -o obj/activations.o
./src/activations.c: In function 'activate':
./src/activations.c:79:5: warning: enumeration value 'RELU6' not handled in
switch [-Wswitch]
    79 |     switch(a){
       |         ^~~~~
./src/activations.c:79:5: warning: enumeration value 'SWISH' not handled in
switch [-Wswitch]
./src/activations.c:79:5: warning: enumeration value 'MISH' not handled in s
witch [-Wswitch]
./src/activations.c:79:5: warning: enumeration value 'HARD_MISH' not handled
in switch [-Wswitch]
./src/activations.c:79:5: warning: enumeration value 'NORM_CHAN' not handled
in switch [-Wswitch]
./src/activations.c:79:5: warning: enumeration value 'NORM_CHAN_SOFTMAX' no
t handled in switch [-Wswitch]
./src/activations.c:79:5: warning: enumeration value 'NORM_CHAN_SOFTMAX_MAX
VAL' not handled in switch [-Wswitch]
./src/activations.c: In function 'gradient':
./src/activations.c:310:5: warning: enumeration value 'SWISH' not handled in
switch [-Wswitch]
   310 |     switch(a){

```

```

^~~~~~
./src/activations.c:310:5: warning: enumeration value 'MISH' not handled in switch [-Wswitch]
./src/activations.c:310:5: warning: enumeration value 'HARD_MISH' not handle d in switch [-Wswitch]
gcc -linclude/ -l3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-u nknown-pragmas -fPIC -rdynamic -Ofast -c ./src/im2col.c -o obj/im2col.o
gcc -linclude/ -l3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-u nknown-pragmas -fPIC -rdynamic -Ofast -c ./src/col2im.c -o obj/col2im.o
gcc -linclude/ -l3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-u nknown-pragmas -fPIC -rdynamic -Ofast -c ./src/blas.c -o obj/blas.o
./src/blas.c: In function 'backward_shortcut_multilayer_cpu':
./src/blas.c:207:21: warning: unused variable 'out_index' [-Wunused-var iab le]
    207 |             int out_index = id;
        |                 ^~~~~~
./src/blas.c: In function 'find_sim':
./src/blas.c:597:59: warning: format '%d' expects argument of type 'int', b ut argument 2 has type 'size_t' {aka 'long unsigned int'} [-Wformat=]
    597 |         printf(" Error: find_sim(): sim isn't found: i = %d, j = %d, z = %d
Wn", i, j, z);
        |                                     ~^
~
|
|
|
size_t {aka long unsigned int}
|
%ld
./src/blas.c:597:67: warning: format '%d' expects argument of type 'int', b ut argument 3 has type 'size_t' {aka 'long unsigned int'} [-Wformat=]
    597 |         printf(" Error: find_sim(): sim isn't found: i = %d, j = %d, z = %d
Wn", i, j, z);
        |                                     ~^
~
|
|
|
size_t {aka long unsigned int}
|
%ld
./src/blas.c:597:75: warning: format '%d' expects argument of type 'int', b ut argument 4 has type 'size_t' {aka 'long unsigned int'} [-Wformat=]
    597 |         printf(" Error: find_sim(): sim isn't found: i = %d, j = %d, z = %d
Wn", i, j, z);
        |                                     ~^
~
|
|
|
size_t {aka long unsigned int}
|
in
t size_t {aka long unsigned int}
|
%ld
./src/blas.c: In function 'find_P_constrastive':
./src/blas.c:611:68: warning: format '%d' expects argument of type 'int', b ut argument 2 has type 'size_t' {aka 'long unsigned int'} [-Wformat=]
    611 |         printf(" Error: find_P_constrastive(): P isn't found: i = %d, j = %
d, z = %d Wn", i, j, z);
        |                                     ~^
~
|
|
|
size_t {aka long unsigned int}
|
%ld
./src/blas.c:611:76: warning: format '%d' expects argument of type 'int', b ut argument 3 has type 'size_t' {aka 'long unsigned int'} [-Wformat=]
    611 |         printf(" Error: find_P_constrastive(): P isn't found: i = %d, j = %d

```

```

, z = %d Wn", i, j, z);
~
|
|
|
nt      size_t {aka long unsigned int}
|
d
./src/blas.c:611:84: warning: format '%d' expects argument of type 'int', but
argument 4 has type 'size_t' {aka 'long unsigned int'} [-Wformat=]
611 |         printf(" Error: find_P_constrastive(): P isn't found: i = %d, j = %
d, z = %d Wn", i, j, z);
~^
~
|
|
int     size_t {aka long unsigned int}
|
%d
./src/blas.c: In function 'P_constrastive_f' :
./src/blas.c:651:79: warning: format '%d' expects argument of type 'int', but
argument 3 has type 'size_t' {aka 'long unsigned int'} [-Wformat=]
651 |         fprintf(stderr, " Error: in P_constrastive must be i != l, while i =
%d, l = %d Wn", i, l);
~^
~
|
|
int     size_t {aka long unsigned int}
|
%d
./src/blas.c:651:87: warning: format '%d' expects argument of type 'int', but
argument 4 has type 'size_t' {aka 'long unsigned int'} [-Wformat=]
651 |         fprintf(stderr, " Error: in P_constrastive must be i != l, while i =
%d, l = %d Wn", i, l);
~^
~
|
|
int     size_t {aka long unsigned int}
|
%d
./src/blas.c: In function 'P_constrastive' :
./src/blas.c:785:79: warning: format '%d' expects argument of type 'int', but
argument 3 has type 'size_t' {aka 'long unsigned int'} [-Wformat=]
785 |         fprintf(stderr, " Error: in P_constrastive must be i != l, while i =
%d, l = %d Wn", i, l);
~^
~
|
|
int     size_t {aka long unsigned int}
|
%d
./src/blas.c:785:87: warning: format '%d' expects argument of type 'int', but
argument 4 has type 'size_t' {aka 'long unsigned int'} [-Wformat=]
785 |         fprintf(stderr, " Error: in P_constrastive must be i != l, while i =
%d, l = %d Wn", i, l);
~^
~

```

```

|
|
int      size_t {aka long unsigned int}

%ld
gcc -linclude/ -l3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-unknown-pragmas -fPIC -rdynamic -Ofast -c ./src/crop_layer.c -o obj/crop_layer.o
gcc -linclude/ -l3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-unknown-pragmas -fPIC -rdynamic -Ofast -c ./src/dropout_layer.c -o obj/dropout_layer.o
gcc -linclude/ -l3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-unknown-pragmas -fPIC -rdynamic -Ofast -c ./src/maxpool_layer.c -o obj/maxpool_layer.o
gcc -linclude/ -l3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-unknown-pragmas -fPIC -rdynamic -Ofast -c ./src/softmax_layer.c -o obj/softmax_layer.o

./src/softmax_layer.c: In function 'make_contrastive_layer':
./src/softmax_layer.c:203:101: warning: format '%d' expects argument of type 'int', but argument 9 has type 'size_t' {aka 'const long unsigned int'} [-Wformat=]
203 |         fprintf(stderr, "contrastive %4d x%4d x%4d x emb_size %4d x batch: %4d classes = %4d, step = %4d Wn", w, h, l.n, l.embedding_size, batch, l.classes, step);
    |                                                                                                     ~~~~
    |                                                                                                     |
int                                                                                                     size_t {aka const long unsigned int}
ed int}

%ld
./src/softmax_layer.c: In function 'forward_contrastive_layer':
./src/softmax_layer.c:244:27: warning: variable 'max_truth' set but not used [-Wunused-but-set-variable]
244 |         float max_truth = 0;
    |                ^~~~~~

gcc -linclude/ -l3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-unknown-pragmas -fPIC -rdynamic -Ofast -c ./src/data.c -o obj/data.o
./src/data.c: In function 'load_data_detection':
./src/data.c:1409:43: warning: variable 'r_scale' set but not used [-Wunused-but-set-variable]
1409 |         float r1 = 0, r2 = 0, r3 = 0, r4 = 0, r_scale;
    |                                           ^~~~~~

./src/data.c: In function 'fill_truth_detection':
./src/data.c:440:33: warning: '%s' directive writing up to 4095 bytes into a region of size 251 [-Wformat-overflow=]
440 |         sprintf(buff, "echo %s WWrong annotation: w = %fW" >> bad_label.list", labelpath, w);
    |                                ^~
~~~~~
./src/data.c:440:27: note: assuming directive output of 8 bytes
440 |         sprintf(buff, "echo %s WWrong annotation: w = %fW" >> bad_label.list", labelpath, w);
    |                                ^~~~~~
~~~~~
In file included from /usr/include/stdio.h:867,
                 from include/darknet.h:13,
                 from ./src/data.h:5,
                 from ./src/data.c:1:
/usr/include/x86_64-linux-gnu/bits/stdio2.h:36:10: note: '__builtin___sprintf_chk' output between 52 and 4461 bytes into a destination of size 256
36 |         return __builtin___sprintf_chk(__s, __USE_FORTIFY_LEVEL - 1,
    |                ^~~~~~

```

```

37 |         __bos (__s), __fmt, __va_arg_pack ());
    |         ~~~~~
./src/data.c:447:33: warning: 's' directive writing up to 4095 bytes into a
region of size 251 [-Wformat-overflow=]
447 |         sprintf(buff, "echo %s W"Wrong annotation: h = %fW" >> bad_labe
l.list", labelpath, h);
    |                                     ^~
~~~~~
./src/data.c:447:27: note: assuming directive output of 8 bytes
447 |         sprintf(buff, "echo %s W"Wrong annotation: h = %fW" >> bad
_label.list", labelpath, h);
    |                                     ^~~~~~
~~~~~
In file included from /usr/include/stdio.h:867,
from include/darknet.h:13,
from ./src/data.h:5,
from ./src/data.c:1:
/usr/include/x86_64-linux-gnu/bits/stdio2.h:36:10: note: '__builtin__spr
intf_chk' output between 52 and 4461 bytes into a destination of size 256
36 |     return __builtin__sprintf_chk (__s, __USE_FORTIFY_LEVEL - 1,
    |             ^~~~~~
37 |         __bos (__s), __fmt, __va_arg_pack ());
    |         ~~~~~
./src/data.c:432:33: warning: 's' directive writing up to 4095 bytes into a
region of size 251 [-Wformat-overflow=]
432 |         sprintf(buff, "echo %s W"Wrong annotation: x = %f, y = %fW" >> b
ad_label.list", labelpath, x, y);
    |                                     ^~
~~~~~
./src/data.c:432:27: note: assuming directive output of 8 bytes
432 |         sprintf(buff, "echo %s W"Wrong annotation: x = %f, y = %f
W" >> bad_label.list", labelpath, x, y);
    |                                     ^~~~~~
~~~~~
./src/data.c:432:27: note: assuming directive output of 8 bytes
In file included from /usr/include/stdio.h:867,
from include/darknet.h:13,
from ./src/data.h:5,
from ./src/data.c:1:
/usr/include/x86_64-linux-gnu/bits/stdio2.h:36:10: note: '__builtin__spr
intf_chk' output between 61 and 4784 bytes into a destination of size 256
36 |     return __builtin__sprintf_chk (__s, __USE_FORTIFY_LEVEL - 1,
    |             ^~~~~~
37 |         __bos (__s), __fmt, __va_arg_pack ());
    |         ~~~~~
./src/data.c:424:33: warning: 's' directive writing up to 4095 bytes into a
region of size 251 [-Wformat-overflow=]
424 |         sprintf(buff, "echo %s W"Wrong annotation: x = 0 or y = 0W" >> b
ad_label.list", labelpath);
    |                                     ^~
~~~~~
In file included from /usr/include/stdio.h:867,
from include/darknet.h:13,
from ./src/data.h:5,
from ./src/data.c:1:
/usr/include/x86_64-linux-gnu/bits/stdio2.h:36:10: note: '__builtin__spr
intf_chk' output between 59 and 4154 bytes into a destination of size 256
36 |     return __builtin__sprintf_chk (__s, __USE_FORTIFY_LEVEL - 1,
    |             ^~~~~~
37 |         __bos (__s), __fmt, __va_arg_pack ());
    |         ~~~~~
./src/data.c:410:33: warning: 's' directive writing up to 4095 bytes into a
region of size 251 [-Wformat-overflow=]
410 |         sprintf(buff, "echo %s W"Wrong annotation: class_id = %d. But cl
ass_id should be [from 0 to %d]W" >> bad_label.list", labelpath, id, (classes-1));

```

```

|                                     ^~
~~~~~
In file included from /usr/include/stdio.h:867,
                 from include/darknet.h:13,
                 from ./src/data.h:5,
                 from ./src/data.c:1:
/usr/include/x86_64-linux-gnu/bits/stdio2.h:36:10: note: '__builtin___spr
intf_chk' output between 95 and 4210 bytes into a destination of size 256
36 |     return __builtin___sprintf_chk (__s, __USE_FORTIFY_LEVEL - 1,
    |           ^~~~~~
37 |         __bos (__s), __fmt, __va_arg_pack ());
    |           ~~~~~
gcc -linclude/ -l3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-u
nknown-pragmas -fPIC -rdynamic -Ofast -c ./src/matrix.c -o obj/matrix.o
gcc -linclude/ -l3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-u
nknown-pragmas -fPIC -rdynamic -Ofast -c ./src/network.c -o obj/network.o
./src/network.c: In function 'train_network_waitkey':
./src/network.c:435:13: warning: unused variable 'ema_period' [-Wunused-va
riable]
435 |         int ema_period = (net.max_batches - ema_start_point - 1000) * (1.0
- net.ema_alpha);
    |         ^~~~~~
At top level:
./src/network.c:1269:14: warning: 'relu' defined but not used [-Wunused-fun
ction]
1269 | static float relu(float src) {
    |           ^~~~
gcc -linclude/ -l3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-u
nknown-pragmas -fPIC -rdynamic -Ofast -c ./src/connected_layer.c -o obj/connected_la
yer.o
gcc -linclude/ -l3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-u
nknown-pragmas -fPIC -rdynamic -Ofast -c ./src/cost_layer.c -o obj/cost_layer.o
gcc -linclude/ -l3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-u
nknown-pragmas -fPIC -rdynamic -Ofast -c ./src/parser.c -o obj/parser.o
./src/parser.c: In function 'save_implicit_weights':
./src/parser.c:1909:9: warning: unused variable 'i' [-Wunused-variable]
1909 |     int i;
    |     ^
./src/parser.c: In function 'get_classes_multipliers':
./src/parser.c:438:40: warning: argument 1 range [18446744071562067968, 1844674
4073709551615] exceeds maximum object size 9223372036854775807 [-Walloc-size-larg
er-than=]
438 |         classes_multipliers = (float *)calloc(classes_counters, sizeof
(float));
    |                                           ^~~~~~
~~~~~
In file included from ./src/parser.c:3:
/usr/include/stdlib.h:542:14: note: in a call to allocation function 'calloc
' declared here
542 | extern void *calloc (size_t __nmemb, size_t __size)
    |           ^~~~~~
gcc -linclude/ -l3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-u
nknown-pragmas -fPIC -rdynamic -Ofast -c ./src/option_list.c -o obj/option_list.o
gcc -linclude/ -l3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-u
nknown-pragmas -fPIC -rdynamic -Ofast -c ./src/darknet.c -o obj/darknet.o
gcc -linclude/ -l3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-u
nknown-pragmas -fPIC -rdynamic -Ofast -c ./src/detection_layer.c -o obj/detection_la
yer.o
gcc -linclude/ -l3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-u
nknown-pragmas -fPIC -rdynamic -Ofast -c ./src/captcha.c -o obj/captcha.o
gcc -linclude/ -l3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-u
nknown-pragmas -fPIC -rdynamic -Ofast -c ./src/route_layer.c -o obj/route_layer.o
gcc -linclude/ -l3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-u
nknown-pragmas -fPIC -rdynamic -Ofast -c ./src/writing.c -o obj/writing.o
gcc -linclude/ -l3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-u

```



```

nknown-pragmas -fPIC -rdynamic -Ofast -c ./src/box.c -o obj/box.o
./src/box.c: In function 'box_iou_kind' :
./src/box.c:154:5: warning: enumeration value 'MSE' not handled in switch [-W
switch]
154 |         switch(iou_kind) {
    |         ^~~~~~
./src/box.c: In function 'diounms_sort' :
./src/box.c:898:27: warning: unused variable 'beta_prob' [-Wunused-variabl
e]
898 |         float beta_prob = pow(dets[j].prob[k], 2) / sum_prob;
    |         ^~~~~~
./src/box.c:897:27: warning: unused variable 'alpha_prob' [-Wunused-variab
le]
897 |         float alpha_prob = pow(dets[i].prob[k], 2) / sum_prob;
    |         ^~~~~~
gcc -linclude/ -l3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-u
nknown-pragmas -fPIC -rdynamic -Ofast -c ./src/nightmare.c -o obj/nightmare.o
gcc -linclude/ -l3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-u
nknown-pragmas -fPIC -rdynamic -Ofast -c ./src/normalization_layer.c -o obj/normaliz
ation_layer.o
gcc -linclude/ -l3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-u
nknown-pragmas -fPIC -rdynamic -Ofast -c ./src/avgpool_layer.c -o obj/avgpool_layer.
o
gcc -linclude/ -l3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-u
nknown-pragmas -fPIC -rdynamic -Ofast -c ./src/coco.c -o obj/coco.o
./src/coco.c: In function 'validate_coco_recall' :
./src/coco.c:248:11: warning: unused variable 'base' [-Wunused-variable]
248 |     char *base = "results/comp4_det_test_";
    |     ^~~~
gcc -linclude/ -l3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-u
nknown-pragmas -fPIC -rdynamic -Ofast -c ./src/dice.c -o obj/dice.o
gcc -linclude/ -l3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-u
nknown-pragmas -fPIC -rdynamic -Ofast -c ./src/yolo.c -o obj/yolo.o
gcc -linclude/ -l3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-u
nknown-pragmas -fPIC -rdynamic -Ofast -c ./src/detector.c -o obj/detector.o
./src/detector.c: In function 'train_detector' :
./src/detector.c:395:72: warning: suggest parentheses around '&&' within '||'
' [-Wparentheses]
395 |         (iteration >= (iter_save + 1000) || iteration % 1000 ==
0) && net.max_batches < 10000)
    |         ~~~~~
~~~~~
./src/detector.c:328:13: warning: variable 'draw_precision' set but not use
d [-Wunused-but-set-variable]
328 |         int draw_precision = 0;
    |         ^~~~~~
./src/detector.c:67:11: warning: unused variable 'avg_contrastive_acc' [-W
unused-variable]
67 |         float avg_contrastive_acc = 0;
    |         ^~~~~~
./src/detector.c: In function 'eliminate_bdd' :
./src/detector.c:588:21: warning: statement with no effect [-Wunused-value]
588 |         for (k; buf[k + n] != 'W0'; k++)
    |         ^~~
./src/detector.c: In function 'validate_detector' :
./src/detector.c:709:13: warning: unused variable 'mkd2' [-Wunused-variabl
e]
709 |         int mkd2 = make_directory(buff2, 0777);
    |         ^~~~
./src/detector.c:707:13: warning: unused variable 'mkd' [-Wunused-variable]
707 |         int mkd = make_directory(buff, 0777);
    |         ^~~
./src/detector.c: In function 'validate_detector_map' :
./src/detector.c:1326:24: warning: variable 'cur_prob' set but not used [-W

```

```

unused-but-set-variable]
1326 |             double cur_prob = 0;
      |             ^~~~~~
./src/detector.c:1347:15: warning: unused variable 'class_recall' [-Wunused-variable]
1347 |             float class_recall = (float)tp_for_thresh_per_class[i] / ((float)tp_for_thresh_per_class[i] + (float)(truth_classes_count[i] - tp_for_thresh_per_class[i]));
      |             ^~~~~~
./src/detector.c:1346:15: warning: unused variable 'class_precision' [-Wunused-variable]
1346 |             float class_precision = (float)tp_for_thresh_per_class[i] / ((float)tp_for_thresh_per_class[i] + (float)fp_for_thresh_per_class[i]);
      |             ^~~~~~
At top level:
./src/detector.c:461:12: warning: 'get_coco_image_id' defined but not used [-Wunused-function]
461 | static int get_coco_image_id(char *filename)
      |             ^~~~~~
gcc -linclude/ -I3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-unknown-pragmas -fPIC -rdynamic -Ofast -c ./src/layer.c -o obj/layer.o
gcc -linclude/ -I3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-unknown-pragmas -fPIC -rdynamic -Ofast -c ./src/compare.c -o obj/compare.o
gcc -linclude/ -I3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-unknown-pragmas -fPIC -rdynamic -Ofast -c ./src/classifier.c -o obj/classifier.o
./src/classifier.c: In function 'train_classifier':
./src/classifier.c:190:13: warning: variable 'draw_precision' set but not used [-Wunused-but-set-variable]
190 |             int draw_precision = 0;
      |             ^~~~~~
./src/classifier.c:146:9: warning: unused variable 'count' [-Wunused-variable]
146 |             int count = 0;
      |             ^~~~~
./src/classifier.c:35:11: warning: unused variable 'avg_contrastive_acc' [-Wunused-variable]
35 |             float avg_contrastive_acc = 0;
      |             ^~~~~~
./src/classifier.c: In function 'predict_classifier':
./src/classifier.c:855:13: warning: unused variable 'time' [-Wunused-variable]
855 |             clock_t time;
      |             ^~~~
gcc -linclude/ -I3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-unknown-pragmas -fPIC -rdynamic -Ofast -c ./src/local_layer.c -o obj/local_layer.o
gcc -linclude/ -I3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-unknown-pragmas -fPIC -rdynamic -Ofast -c ./src/swag.c -o obj/swag.o
gcc -linclude/ -I3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-unknown-pragmas -fPIC -rdynamic -Ofast -c ./src/shortcut_layer.c -o obj/shortcut_layer.o
./src/shortcut_layer.c: In function 'make_shortcut_layer':
./src/shortcut_layer.c:55:15: warning: unused variable 'scale' [-Wunused-variable]
55 |             float scale = sqrt(2. / l.nweights);
      |             ^~~~~
gcc -linclude/ -I3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-unknown-pragmas -fPIC -rdynamic -Ofast -c ./src/representation_layer.c -o obj/representation_layer.o
gcc -linclude/ -I3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-unknown-pragmas -fPIC -rdynamic -Ofast -c ./src/activation_layer.c -o obj/activation_layer.o
gcc -linclude/ -I3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-unknown-pragmas -fPIC -rdynamic -Ofast -c ./src/rnn_layer.c -o obj/rnn_layer.o
gcc -linclude/ -I3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-unknown-pragmas -fPIC -rdynamic -Ofast -c ./src/gru_layer.c -o obj/gru_layer.o

```

```

gcc -linclude/ -l3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-unknown-pragmas -fPIC -rdynamic -Ofast -c ./src/rnn.c -o obj/rnn.o
gcc -linclude/ -l3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-unknown-pragmas -fPIC -rdynamic -Ofast -c ./src/rnn_vid.c -o obj/rnn_vid.o
gcc -linclude/ -l3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-unknown-pragmas -fPIC -rdynamic -Ofast -c ./src/crnn_layer.c -o obj/crnn_layer.o
gcc -linclude/ -l3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-unknown-pragmas -fPIC -rdynamic -Ofast -c ./src/demo.c -o obj/demo.o
gcc -linclude/ -l3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-unknown-pragmas -fPIC -rdynamic -Ofast -c ./src/tag.c -o obj/tag.o
gcc -linclude/ -l3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-unknown-pragmas -fPIC -rdynamic -Ofast -c ./src/cifar.c -o obj/cifar.o
gcc -linclude/ -l3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-unknown-pragmas -fPIC -rdynamic -Ofast -c ./src/go.c -o obj/go.o
gcc -linclude/ -l3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-unknown-pragmas -fPIC -rdynamic -Ofast -c ./src/batchnorm_layer.c -o obj/batchnorm_layer.o
gcc -linclude/ -l3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-unknown-pragmas -fPIC -rdynamic -Ofast -c ./src/art.c -o obj/art.o
gcc -linclude/ -l3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-unknown-pragmas -fPIC -rdynamic -Ofast -c ./src/region_layer.c -o obj/region_layer.o
gcc -linclude/ -l3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-unknown-pragmas -fPIC -rdynamic -Ofast -c ./src/reorg_layer.c -o obj/reorg_layer.o
gcc -linclude/ -l3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-unknown-pragmas -fPIC -rdynamic -Ofast -c ./src/reorg_old_layer.c -o obj/reorg_old_layer.o
gcc -linclude/ -l3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-unknown-pragmas -fPIC -rdynamic -Ofast -c ./src/super.c -o obj/super.o
gcc -linclude/ -l3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-unknown-pragmas -fPIC -rdynamic -Ofast -c ./src/voxel.c -o obj/voxel.o
gcc -linclude/ -l3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-unknown-pragmas -fPIC -rdynamic -Ofast -c ./src/tree.c -o obj/tree.o
gcc -linclude/ -l3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-unknown-pragmas -fPIC -rdynamic -Ofast -c ./src/yolo_layer.c -o obj/yolo_layer.o
./src/yolo_layer.c: In function 'process_batch':
./src/yolo_layer.c:426:25: warning: variable 'best_match_t' set but not used [-Wuninitialized]
    426 |         int best_match_t = 0;
        |         ^~~~~~
./src/yolo_layer.c: In function 'forward_yolo_layer':
./src/yolo_layer.c:707:11: warning: unused variable 'avg_anyobj' [-Wunused-variable]
    707 |     float avg_anyobj = 0;
        |     ^~~~~~
./src/yolo_layer.c:706:11: warning: unused variable 'avg_obj' [-Wunused-variable]
    706 |     float avg_obj = 0;
        |     ^~~~~~
./src/yolo_layer.c:705:11: warning: unused variable 'avg_cat' [-Wunused-variable]
    705 |     float avg_cat = 0;
        |     ^~~~~~
./src/yolo_layer.c:704:11: warning: unused variable 'recall75' [-Wunused-variable]
    704 |     float recall75 = 0;
        |     ^~~~~~
./src/yolo_layer.c:703:11: warning: unused variable 'recall' [-Wunused-variable]
    703 |     float recall = 0;
        |     ^~~~~~
./src/yolo_layer.c:702:11: warning: unused variable 'tot_ciou_loss' [-Wunused-variable]
    702 |     float tot_ciou_loss = 0;
        |     ^~~~~~
./src/yolo_layer.c:701:11: warning: unused variable 'tot_diou_loss' [-Wunused-variable]
    701 |     float tot_diou_loss = 0;
        |     ^~~~~~

```

```

sed-variable]
701 |         float tot_diou_loss = 0;
      |         ^~~~~~
./src/yolo_layer.c:698:11: warning: unused variable 'tot_ciou' [-Wunused-variable]
698 |         float tot_ciou = 0;
      |         ^~~~~~
./src/yolo_layer.c:697:11: warning: unused variable 'tot_diou' [-Wunused-variable]
697 |         float tot_diou = 0;
      |         ^~~~~~
./src/yolo_layer.c:696:11: warning: unused variable 'tot_giou' [-Wunused-variable]
696 |         float tot_giou = 0;
      |         ^~~~~~
gcc -linclude/ -l3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-unknown-pragmas -fPIC -rdynamic -Ofast -c ./src/gaussian_yolo_layer.c -o obj/gaussian_yolo_layer.o
gcc -linclude/ -l3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-unknown-pragmas -fPIC -rdynamic -Ofast -c ./src/upsample_layer.c -o obj/upsample_layer.o
gcc -linclude/ -l3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-unknown-pragmas -fPIC -rdynamic -Ofast -c ./src/lstm_layer.c -o obj/lstm_layer.o
gcc -linclude/ -l3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-unknown-pragmas -fPIC -rdynamic -Ofast -c ./src/conv_lstm_layer.c -o obj/conv_lstm_layer.o
gcc -linclude/ -l3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-unknown-pragmas -fPIC -rdynamic -Ofast -c ./src/scale_channels_layer.c -o obj/scale_channels_layer.o
gcc -linclude/ -l3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-unknown-pragmas -fPIC -rdynamic -Ofast -c ./src/sam_layer.c -o obj/sam_layer.o
g++ -std=c++11 -std=c++11 -linclude/ -l3rdparty/stb/include -Wall -Wfatal-errors -Wno-unused-result -Wno-unknown-pragmas -fPIC -rdynamic -Ofast obj/image_opencv.o obj/http_stream.o obj/gemm.o obj/Utils.o obj/dark_cuda.o obj/convolutional_layer.o obj/list.o obj/image.o obj/activations.o obj/im2col.o obj/col2im.o obj/blas.o obj/crop_layer.o obj/dropout_layer.o obj/maxpool_layer.o obj/softmax_layer.o obj/data.o obj/matrix.o obj/network.o obj/connected_layer.o obj/cost_layer.o obj/parser.o obj/option_list.o obj/darknet.o obj/detection_layer.o obj/captcha.o obj/route_layer.o obj/writing.o obj/box.o obj/nightmare.o obj/normalization_layer.o obj/avgpool_layer.o obj/coco.o obj/dice.o obj/yolo.o obj/detector.o obj/layer.o obj/compare.o obj/classifier.o obj/local_layer.o obj/swag.o obj/shortcut_layer.o obj/representation_layer.o obj/activation_layer.o obj/rnn_layer.o obj/gru_layer.o obj/rnn.o obj/rnn_vid.o obj/crnn_layer.o obj/demo.o obj/tag.o obj/cifar.o obj/go.o obj/batchnorm_layer.o obj/art.o obj/region_layer.o obj/reorg_layer.o obj/reorg_old_layer.o obj/super.o obj/voxel.o obj/tree.o obj/yolo_layer.o obj/gaussian_yolo_layer.o obj/upsample_layer.o obj/lstm_layer.o obj/conv_lstm_layer.o obj/scale_channels_layer.o obj/sam_layer.o -o darknet -lm -pthread

```

사전 훈련된 가중치 다운로드

```
In [12]: !cd darknet; wget https://pjreddie.com/media/files/yolov3.weights
```

```

--2023-05-25 13:17:56-- https://pjreddie.com/media/files/yolov3.weights
Resolving pjreddie.com (pjreddie.com)... 128.208.4.108
Connecting to pjreddie.com (pjreddie.com)|128.208.4.108|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 248007048 (237M) [application/octet-stream]
Saving to: 'yolov3.weights'

```

```
yolov3.weights      100%[=====>] 236.52M  20.1MB/s   in 12s
```

```
2023-05-25 13:18:09 (19.3 MB/s) - 'yolov3.weights' saved [248007048/248007048]
```

02 yolo를 활용 - 예측수행

목차로 이동하기

```
In [17]: ### 이미지 확인
import matplotlib.pyplot as plt

# 이미지 파일의 경로를 지정합니다.
image_path = "/content/darknet/data/dog.jpg"

# 이미지를 읽어옵니다.
image = plt.imread(image_path)

# 이미지를 출력합니다.
plt.imshow(image)
plt.show()
```



```
In [20]: !cd darknet; ./darknet detect cfg/yolov3.cfg yolov3.weights data/dog.jpg
```

GPU isn't used
OpenCV isn't used - data augmentation will be slow
mini_batch = 1, batch = 1, time_steps = 1, train = 0

layer	filters	size/strd(dil)	input	output
0 conv	32	3 x 3/ 1	416 x 416 x 3 ->	416 x 416 x 32 0.299 BF
1 conv	64	3 x 3/ 2	416 x 416 x 32 ->	208 x 208 x 64 1.595 BF
2 conv	32	1 x 1/ 1	208 x 208 x 64 ->	208 x 208 x 32 0.177 BF
3 conv	64	3 x 3/ 1	208 x 208 x 32 ->	208 x 208 x 64 1.595 BF
4 Shortcut Layer:	1,	wt = 0, wn = 0,	outputs: 208 x 208 x	64 0.003 BF
5 conv	128	3 x 3/ 2	208 x 208 x 64 ->	104 x 104 x 128 1.595 BF
6 conv	64	1 x 1/ 1	104 x 104 x 128 ->	104 x 104 x 64 0.177 BF
7 conv	128	3 x 3/ 1	104 x 104 x 64 ->	104 x 104 x 128 1.595 BF
8 Shortcut Layer:	5,	wt = 0, wn = 0,	outputs: 104 x 104 x	128 0.001 BF
9 conv	64	1 x 1/ 1	104 x 104 x 128 ->	104 x 104 x 64 0.177 BF
10 conv	128	3 x 3/ 1	104 x 104 x 64 ->	104 x 104 x 128 1.595 BF
11 Shortcut Layer:	8,	wt = 0, wn = 0,	outputs: 104 x 104 x	128 0.001 BF
12 conv	256	3 x 3/ 2	104 x 104 x 128 ->	52 x 52 x 256 1.595 BF
13 conv	128	1 x 1/ 1	52 x 52 x 256 ->	52 x 52 x 128 0.177 BF
14 conv	256	3 x 3/ 1	52 x 52 x 128 ->	52 x 52 x 256 1.595 BF
15 Shortcut Layer:	12,	wt = 0, wn = 0,	outputs: 52 x 52 x	256 0.001 BF
16 conv	128	1 x 1/ 1	52 x 52 x 256 ->	52 x 52 x 128 0.177 BF
17 conv	256	3 x 3/ 1	52 x 52 x 128 ->	52 x 52 x 256 1.595 BF
18 Shortcut Layer:	15,	wt = 0, wn = 0,	outputs: 52 x 52 x	256 0.001 BF
19 conv	128	1 x 1/ 1	52 x 52 x 256 ->	52 x 52 x 128 0.177 BF
20 conv	256	3 x 3/ 1	52 x 52 x 128 ->	52 x 52 x 256 1.595 BF
21 Shortcut Layer:	18,	wt = 0, wn = 0,	outputs: 52 x 52 x	256 0.001 BF
22 conv	128	1 x 1/ 1	52 x 52 x 256 ->	52 x 52 x 128 0.177 BF
23 conv	256	3 x 3/ 1	52 x 52 x 128 ->	52 x 52 x 256 1.595 BF
24 Shortcut Layer:	21,	wt = 0, wn = 0,	outputs: 52 x 52 x	256 0.001 BF
25 conv	128	1 x 1/ 1	52 x 52 x 256 ->	52 x 52 x 128 0.177 BF
26 conv	256	3 x 3/ 1	52 x 52 x 128 ->	52 x 52 x 256 1.595 BF
27 Shortcut Layer:	24,	wt = 0, wn = 0,	outputs: 52 x 52 x	256 0.001 BF
28 conv	128	1 x 1/ 1	52 x 52 x 256 ->	52 x 52 x 128 0.177 BF
29 conv	256	3 x 3/ 1	52 x 52 x 128 ->	52 x 52 x 256 1.595 BF
30 Shortcut Layer:	27,	wt = 0, wn = 0,	outputs: 52 x 52 x	256 0.001 BF
31 conv	128	1 x 1/ 1	52 x 52 x 256 ->	52 x 52 x 128 0.177 BF
32 conv	256	3 x 3/ 1	52 x 52 x 128 ->	52 x 52 x 256 1.595 BF
33 Shortcut Layer:	30,	wt = 0, wn = 0,	outputs: 52 x 52 x	256 0.001 BF
34 conv	128	1 x 1/ 1	52 x 52 x 256 ->	52 x 52 x 128 0.177 BF
35 conv	256	3 x 3/ 1	52 x 52 x 128 ->	52 x 52 x 256 1.595 BF
36 Shortcut Layer:	33,	wt = 0, wn = 0,	outputs: 52 x 52 x	256 0.001 BF
37 conv	512	3 x 3/ 2	52 x 52 x 256 ->	26 x 26 x 512 1.595 BF
38 conv	256	1 x 1/ 1	26 x 26 x 512 ->	26 x 26 x 256 0.177 BF
39 conv	512	3 x 3/ 1	26 x 26 x 256 ->	26 x 26 x 512 1.595 BF
40 Shortcut Layer:	37,	wt = 0, wn = 0,	outputs: 26 x 26 x	512 0.000 BF
41 conv	256	1 x 1/ 1	26 x 26 x 512 ->	26 x 26 x 256 0.177 BF
42 conv	512	3 x 3/ 1	26 x 26 x 256 ->	26 x 26 x 512 1.595 BF
43 Shortcut Layer:	40,	wt = 0, wn = 0,	outputs: 26 x 26 x	512 0.000 BF
44 conv	256	1 x 1/ 1	26 x 26 x 512 ->	26 x 26 x 256 0.177 BF
45 conv	512	3 x 3/ 1	26 x 26 x 256 ->	26 x 26 x 512 1.595 BF
46 Shortcut Layer:	43,	wt = 0, wn = 0,	outputs: 26 x 26 x	512 0.000 BF
47 conv	256	1 x 1/ 1	26 x 26 x 512 ->	26 x 26 x 256 0.177 BF
48 conv	512	3 x 3/ 1	26 x 26 x 256 ->	26 x 26 x 512 1.595 BF
49 Shortcut Layer:	46,	wt = 0, wn = 0,	outputs: 26 x 26 x	512 0.000 BF
50 conv	256	1 x 1/ 1	26 x 26 x 512 ->	26 x 26 x 256 0.177 BF
51 conv	512	3 x 3/ 1	26 x 26 x 256 ->	26 x 26 x 512 1.595 BF
52 Shortcut Layer:	49,	wt = 0, wn = 0,	outputs: 26 x 26 x	512 0.000 BF
53 conv	256	1 x 1/ 1	26 x 26 x 512 ->	26 x 26 x 256 0.177 BF
54 conv	512	3 x 3/ 1	26 x 26 x 256 ->	26 x 26 x 512 1.595 BF
55 Shortcut Layer:	52,	wt = 0, wn = 0,	outputs: 26 x 26 x	512 0.000 BF
56 conv	256	1 x 1/ 1	26 x 26 x 512 ->	26 x 26 x 256 0.177 BF
57 conv	512	3 x 3/ 1	26 x 26 x 256 ->	26 x 26 x 512 1.595 BF
58 Shortcut Layer:	55,	wt = 0, wn = 0,	outputs: 26 x 26 x	512 0.000 BF
59 conv	256	1 x 1/ 1	26 x 26 x 512 ->	26 x 26 x 256 0.177 BF
60 conv	512	3 x 3/ 1	26 x 26 x 256 ->	26 x 26 x 512 1.595 BF

```

61 Shortcut Layer: 58, wt = 0, wn = 0, outputs: 26 x 26 x 512 0.000 BF
62 conv 1024 3 x 3/ 2 26 x 26 x 512 -> 13 x 13 x 1024 1.595 BF
63 conv 512 1 x 1/ 1 13 x 13 x 1024 -> 13 x 13 x 512 0.177 BF
64 conv 1024 3 x 3/ 1 13 x 13 x 512 -> 13 x 13 x 1024 1.595 BF
65 Shortcut Layer: 62, wt = 0, wn = 0, outputs: 13 x 13 x 1024 0.000 BF
66 conv 512 1 x 1/ 1 13 x 13 x 1024 -> 13 x 13 x 512 0.177 BF
67 conv 1024 3 x 3/ 1 13 x 13 x 512 -> 13 x 13 x 1024 1.595 BF
68 Shortcut Layer: 65, wt = 0, wn = 0, outputs: 13 x 13 x 1024 0.000 BF
69 conv 512 1 x 1/ 1 13 x 13 x 1024 -> 13 x 13 x 512 0.177 BF
70 conv 1024 3 x 3/ 1 13 x 13 x 512 -> 13 x 13 x 1024 1.595 BF
71 Shortcut Layer: 68, wt = 0, wn = 0, outputs: 13 x 13 x 1024 0.000 BF
72 conv 512 1 x 1/ 1 13 x 13 x 1024 -> 13 x 13 x 512 0.177 BF
73 conv 1024 3 x 3/ 1 13 x 13 x 512 -> 13 x 13 x 1024 1.595 BF
74 Shortcut Layer: 71, wt = 0, wn = 0, outputs: 13 x 13 x 1024 0.000 BF
75 conv 512 1 x 1/ 1 13 x 13 x 1024 -> 13 x 13 x 512 0.177 BF
76 conv 1024 3 x 3/ 1 13 x 13 x 512 -> 13 x 13 x 1024 1.595 BF
77 conv 512 1 x 1/ 1 13 x 13 x 1024 -> 13 x 13 x 512 0.177 BF
78 conv 1024 3 x 3/ 1 13 x 13 x 512 -> 13 x 13 x 1024 1.595 BF
79 conv 512 1 x 1/ 1 13 x 13 x 1024 -> 13 x 13 x 512 0.177 BF
80 conv 1024 3 x 3/ 1 13 x 13 x 512 -> 13 x 13 x 1024 1.595 BF
81 conv 255 1 x 1/ 1 13 x 13 x 1024 -> 13 x 13 x 255 0.088 BF
82 yolo
[yolo] params: iou_loss: mse (2), iou_norm: 0.75, obj_norm: 1.00, cls_norm: 1.00, de
lta_norm: 1.00, scale_x_y: 1.00
83 route 79 -> 13 x 13 x 512
84 conv 256 1 x 1/ 1 13 x 13 x 512 -> 13 x 13 x 256 0.044 BF
85 upsample 2x 13 x 13 x 256 -> 26 x 26 x 256
86 route 85 61 -> 26 x 26 x 768
87 conv 256 1 x 1/ 1 26 x 26 x 768 -> 26 x 26 x 256 0.266 BF
88 conv 512 3 x 3/ 1 26 x 26 x 256 -> 26 x 26 x 512 1.595 BF
89 conv 256 1 x 1/ 1 26 x 26 x 512 -> 26 x 26 x 256 0.177 BF
90 conv 512 3 x 3/ 1 26 x 26 x 256 -> 26 x 26 x 512 1.595 BF
91 conv 256 1 x 1/ 1 26 x 26 x 512 -> 26 x 26 x 256 0.177 BF
92 conv 512 3 x 3/ 1 26 x 26 x 256 -> 26 x 26 x 512 1.595 BF
93 conv 255 1 x 1/ 1 26 x 26 x 512 -> 26 x 26 x 255 0.177 BF
94 yolo
[yolo] params: iou_loss: mse (2), iou_norm: 0.75, obj_norm: 1.00, cls_norm: 1.00, de
lta_norm: 1.00, scale_x_y: 1.00
95 route 91 -> 26 x 26 x 256
96 conv 128 1 x 1/ 1 26 x 26 x 256 -> 26 x 26 x 128 0.044 BF
97 upsample 2x 26 x 26 x 128 -> 52 x 52 x 128
98 route 97 36 -> 52 x 52 x 384
99 conv 128 1 x 1/ 1 52 x 52 x 384 -> 52 x 52 x 128 0.266 BF
100 conv 256 3 x 3/ 1 52 x 52 x 128 -> 52 x 52 x 256 1.595 BF
101 conv 128 1 x 1/ 1 52 x 52 x 256 -> 52 x 52 x 128 0.177 BF
102 conv 256 3 x 3/ 1 52 x 52 x 128 -> 52 x 52 x 256 1.595 BF
103 conv 128 1 x 1/ 1 52 x 52 x 256 -> 52 x 52 x 128 0.177 BF
104 conv 256 3 x 3/ 1 52 x 52 x 128 -> 52 x 52 x 256 1.595 BF
105 conv 255 1 x 1/ 1 52 x 52 x 256 -> 52 x 52 x 255 0.353 BF
106 yolo
[yolo] params: iou_loss: mse (2), iou_norm: 0.75, obj_norm: 1.00, cls_norm: 1.00, de
lta_norm: 1.00, scale_x_y: 1.00
Total BFLOPS 65.879
avg_outputs = 532444
Loading weights from yolov3.weights...
seen 64, trained: 32013 K-images (500 Kilo-batches_64)
Done! Loaded 107 layers from weights-file
Detection layer: 82 - type = 28
Detection layer: 94 - type = 28
Detection layer: 106 - type = 28
data/dog.jpg: Predicted in 9533.226000 milli-seconds.
bicycle: 99%
dog: 100%
truck: 93%
Not compiled with OpenCV, saving to predictions.png instead

```


실습


- eagle.jpg를 해보기

```
In [21]: ### 이미지 확인
import matplotlib.pyplot as plt

# 이미지 파일의 경로를 지정합니다.
image_path = "/content/darknet/data/eagle.jpg"

# 이미지를 읽어옵니다.
image = plt.imread(image_path)

# 이미지를 출력합니다.
plt.imshow(image)
plt.show()
```



```
In [26]: !cd darknet; ./darknet detect cfg/yolov3.cfg yolov3.weights data/eagle.jpg
```


GPU isn't used
 OpenCV isn't used - data augmentation will be slow
 mini_batch = 1, batch = 1, time_steps = 1, train = 0

layer	filters	size/strd(dil)	input	output
0 conv	32	3 x 3/ 1	416 x 416 x 3 ->	416 x 416 x 32 0.299 BF
1 conv	64	3 x 3/ 2	416 x 416 x 32 ->	208 x 208 x 64 1.595 BF
2 conv	32	1 x 1/ 1	208 x 208 x 64 ->	208 x 208 x 32 0.177 BF
3 conv	64	3 x 3/ 1	208 x 208 x 32 ->	208 x 208 x 64 1.595 BF
4 Shortcut Layer:	1,	wt = 0, wn = 0,	outputs: 208 x 208 x	64 0.003 BF
5 conv	128	3 x 3/ 2	208 x 208 x 64 ->	104 x 104 x 128 1.595 BF
6 conv	64	1 x 1/ 1	104 x 104 x 128 ->	104 x 104 x 64 0.177 BF
7 conv	128	3 x 3/ 1	104 x 104 x 64 ->	104 x 104 x 128 1.595 BF
8 Shortcut Layer:	5,	wt = 0, wn = 0,	outputs: 104 x 104 x	128 0.001 BF
9 conv	64	1 x 1/ 1	104 x 104 x 128 ->	104 x 104 x 64 0.177 BF
10 conv	128	3 x 3/ 1	104 x 104 x 64 ->	104 x 104 x 128 1.595 BF
11 Shortcut Layer:	8,	wt = 0, wn = 0,	outputs: 104 x 104 x	128 0.001 BF
12 conv	256	3 x 3/ 2	104 x 104 x 128 ->	52 x 52 x 256 1.595 BF
13 conv	128	1 x 1/ 1	52 x 52 x 256 ->	52 x 52 x 128 0.177 BF
14 conv	256	3 x 3/ 1	52 x 52 x 128 ->	52 x 52 x 256 1.595 BF
15 Shortcut Layer:	12,	wt = 0, wn = 0,	outputs: 52 x 52 x	256 0.001 BF
16 conv	128	1 x 1/ 1	52 x 52 x 256 ->	52 x 52 x 128 0.177 BF
17 conv	256	3 x 3/ 1	52 x 52 x 128 ->	52 x 52 x 256 1.595 BF
18 Shortcut Layer:	15,	wt = 0, wn = 0,	outputs: 52 x 52 x	256 0.001 BF
19 conv	128	1 x 1/ 1	52 x 52 x 256 ->	52 x 52 x 128 0.177 BF
20 conv	256	3 x 3/ 1	52 x 52 x 128 ->	52 x 52 x 256 1.595 BF
21 Shortcut Layer:	18,	wt = 0, wn = 0,	outputs: 52 x 52 x	256 0.001 BF
22 conv	128	1 x 1/ 1	52 x 52 x 256 ->	52 x 52 x 128 0.177 BF
23 conv	256	3 x 3/ 1	52 x 52 x 128 ->	52 x 52 x 256 1.595 BF
24 Shortcut Layer:	21,	wt = 0, wn = 0,	outputs: 52 x 52 x	256 0.001 BF
25 conv	128	1 x 1/ 1	52 x 52 x 256 ->	52 x 52 x 128 0.177 BF
26 conv	256	3 x 3/ 1	52 x 52 x 128 ->	52 x 52 x 256 1.595 BF
27 Shortcut Layer:	24,	wt = 0, wn = 0,	outputs: 52 x 52 x	256 0.001 BF
28 conv	128	1 x 1/ 1	52 x 52 x 256 ->	52 x 52 x 128 0.177 BF
29 conv	256	3 x 3/ 1	52 x 52 x 128 ->	52 x 52 x 256 1.595 BF
30 Shortcut Layer:	27,	wt = 0, wn = 0,	outputs: 52 x 52 x	256 0.001 BF
31 conv	128	1 x 1/ 1	52 x 52 x 256 ->	52 x 52 x 128 0.177 BF
32 conv	256	3 x 3/ 1	52 x 52 x 128 ->	52 x 52 x 256 1.595 BF
33 Shortcut Layer:	30,	wt = 0, wn = 0,	outputs: 52 x 52 x	256 0.001 BF
34 conv	128	1 x 1/ 1	52 x 52 x 256 ->	52 x 52 x 128 0.177 BF
35 conv	256	3 x 3/ 1	52 x 52 x 128 ->	52 x 52 x 256 1.595 BF
36 Shortcut Layer:	33,	wt = 0, wn = 0,	outputs: 52 x 52 x	256 0.001 BF
37 conv	512	3 x 3/ 2	52 x 52 x 256 ->	26 x 26 x 512 1.595 BF
38 conv	256	1 x 1/ 1	26 x 26 x 512 ->	26 x 26 x 256 0.177 BF
39 conv	512	3 x 3/ 1	26 x 26 x 256 ->	26 x 26 x 512 1.595 BF
40 Shortcut Layer:	37,	wt = 0, wn = 0,	outputs: 26 x 26 x	512 0.000 BF
41 conv	256	1 x 1/ 1	26 x 26 x 512 ->	26 x 26 x 256 0.177 BF
42 conv	512	3 x 3/ 1	26 x 26 x 256 ->	26 x 26 x 512 1.595 BF
43 Shortcut Layer:	40,	wt = 0, wn = 0,	outputs: 26 x 26 x	512 0.000 BF
44 conv	256	1 x 1/ 1	26 x 26 x 512 ->	26 x 26 x 256 0.177 BF
45 conv	512	3 x 3/ 1	26 x 26 x 256 ->	26 x 26 x 512 1.595 BF
46 Shortcut Layer:	43,	wt = 0, wn = 0,	outputs: 26 x 26 x	512 0.000 BF
47 conv	256	1 x 1/ 1	26 x 26 x 512 ->	26 x 26 x 256 0.177 BF
48 conv	512	3 x 3/ 1	26 x 26 x 256 ->	26 x 26 x 512 1.595 BF
49 Shortcut Layer:	46,	wt = 0, wn = 0,	outputs: 26 x 26 x	512 0.000 BF
50 conv	256	1 x 1/ 1	26 x 26 x 512 ->	26 x 26 x 256 0.177 BF
51 conv	512	3 x 3/ 1	26 x 26 x 256 ->	26 x 26 x 512 1.595 BF
52 Shortcut Layer:	49,	wt = 0, wn = 0,	outputs: 26 x 26 x	512 0.000 BF
53 conv	256	1 x 1/ 1	26 x 26 x 512 ->	26 x 26 x 256 0.177 BF
54 conv	512	3 x 3/ 1	26 x 26 x 256 ->	26 x 26 x 512 1.595 BF
55 Shortcut Layer:	52,	wt = 0, wn = 0,	outputs: 26 x 26 x	512 0.000 BF
56 conv	256	1 x 1/ 1	26 x 26 x 512 ->	26 x 26 x 256 0.177 BF
57 conv	512	3 x 3/ 1	26 x 26 x 256 ->	26 x 26 x 512 1.595 BF
58 Shortcut Layer:	55,	wt = 0, wn = 0,	outputs: 26 x 26 x	512 0.000 BF
59 conv	256	1 x 1/ 1	26 x 26 x 512 ->	26 x 26 x 256 0.177 BF
60 conv	512	3 x 3/ 1	26 x 26 x 256 ->	26 x 26 x 512 1.595 BF

```

61 Shortcut Layer: 58, wt = 0, wn = 0, outputs: 26 x 26 x 512 0.000 BF
62 conv 1024 3 x 3/ 2 26 x 26 x 512 -> 13 x 13 x1024 1.595 BF
63 conv 512 1 x 1/ 1 13 x 13 x1024 -> 13 x 13 x 512 0.177 BF
64 conv 1024 3 x 3/ 1 13 x 13 x 512 -> 13 x 13 x1024 1.595 BF
65 Shortcut Layer: 62, wt = 0, wn = 0, outputs: 13 x 13 x1024 0.000 BF
66 conv 512 1 x 1/ 1 13 x 13 x1024 -> 13 x 13 x 512 0.177 BF
67 conv 1024 3 x 3/ 1 13 x 13 x 512 -> 13 x 13 x1024 1.595 BF
68 Shortcut Layer: 65, wt = 0, wn = 0, outputs: 13 x 13 x1024 0.000 BF
69 conv 512 1 x 1/ 1 13 x 13 x1024 -> 13 x 13 x 512 0.177 BF
70 conv 1024 3 x 3/ 1 13 x 13 x 512 -> 13 x 13 x1024 1.595 BF
71 Shortcut Layer: 68, wt = 0, wn = 0, outputs: 13 x 13 x1024 0.000 BF
72 conv 512 1 x 1/ 1 13 x 13 x1024 -> 13 x 13 x 512 0.177 BF
73 conv 1024 3 x 3/ 1 13 x 13 x 512 -> 13 x 13 x1024 1.595 BF
74 Shortcut Layer: 71, wt = 0, wn = 0, outputs: 13 x 13 x1024 0.000 BF
75 conv 512 1 x 1/ 1 13 x 13 x1024 -> 13 x 13 x 512 0.177 BF
76 conv 1024 3 x 3/ 1 13 x 13 x 512 -> 13 x 13 x1024 1.595 BF
77 conv 512 1 x 1/ 1 13 x 13 x1024 -> 13 x 13 x 512 0.177 BF
78 conv 1024 3 x 3/ 1 13 x 13 x 512 -> 13 x 13 x1024 1.595 BF
79 conv 512 1 x 1/ 1 13 x 13 x1024 -> 13 x 13 x 512 0.177 BF
80 conv 1024 3 x 3/ 1 13 x 13 x 512 -> 13 x 13 x1024 1.595 BF
81 conv 255 1 x 1/ 1 13 x 13 x1024 -> 13 x 13 x 255 0.088 BF
82 yolo
[yolo] params: iou_loss: mse (2), iou_norm: 0.75, obj_norm: 1.00, cls_norm: 1.00, de
lta_norm: 1.00, scale_x_y: 1.00
83 route 79 -> 13 x 13 x 512
84 conv 256 1 x 1/ 1 13 x 13 x 512 -> 13 x 13 x 256 0.044 BF
85 upsample 2x 13 x 13 x 256 -> 26 x 26 x 256
86 route 85 61 -> 26 x 26 x 768
87 conv 256 1 x 1/ 1 26 x 26 x 768 -> 26 x 26 x 256 0.266 BF
88 conv 512 3 x 3/ 1 26 x 26 x 256 -> 26 x 26 x 512 1.595 BF
89 conv 256 1 x 1/ 1 26 x 26 x 512 -> 26 x 26 x 256 0.177 BF
90 conv 512 3 x 3/ 1 26 x 26 x 256 -> 26 x 26 x 512 1.595 BF
91 conv 256 1 x 1/ 1 26 x 26 x 512 -> 26 x 26 x 256 0.177 BF
92 conv 512 3 x 3/ 1 26 x 26 x 256 -> 26 x 26 x 512 1.595 BF
93 conv 255 1 x 1/ 1 26 x 26 x 512 -> 26 x 26 x 255 0.177 BF
94 yolo
[yolo] params: iou_loss: mse (2), iou_norm: 0.75, obj_norm: 1.00, cls_norm: 1.00, de
lta_norm: 1.00, scale_x_y: 1.00
95 route 91 -> 26 x 26 x 256
96 conv 128 1 x 1/ 1 26 x 26 x 256 -> 26 x 26 x 128 0.044 BF
97 upsample 2x 26 x 26 x 128 -> 52 x 52 x 128
98 route 97 36 -> 52 x 52 x 384
99 conv 128 1 x 1/ 1 52 x 52 x 384 -> 52 x 52 x 128 0.266 BF
100 conv 256 3 x 3/ 1 52 x 52 x 128 -> 52 x 52 x 256 1.595 BF
101 conv 128 1 x 1/ 1 52 x 52 x 256 -> 52 x 52 x 128 0.177 BF
102 conv 256 3 x 3/ 1 52 x 52 x 128 -> 52 x 52 x 256 1.595 BF
103 conv 128 1 x 1/ 1 52 x 52 x 256 -> 52 x 52 x 128 0.177 BF
104 conv 256 3 x 3/ 1 52 x 52 x 128 -> 52 x 52 x 256 1.595 BF
105 conv 255 1 x 1/ 1 52 x 52 x 256 -> 52 x 52 x 255 0.353 BF
106 yolo
[yolo] params: iou_loss: mse (2), iou_norm: 0.75, obj_norm: 1.00, cls_norm: 1.00, de
lta_norm: 1.00, scale_x_y: 1.00
Total BFLOPS 65.879
avg_outputs = 532444
Loading weights from yolov3.weights...
seen 64, trained: 32013 K-images (500 Kilo-batches_64)
Done! Loaded 107 layers from weights-file
Detection layer: 82 - type = 28
Detection layer: 94 - type = 28
Detection layer: 106 - type = 28
data/eagle.jpg: Predicted in 8527.034000 milli-seconds.
bird: 99%
Not compiled with OpenCV, saving to predictions.png instead

```


03 yolo 예측 결과 확인

목차로 이동하기

```
In [28]: import cv2
import matplotlib.pyplot as plt

def show_detection_result(path):
    image = cv2.imread(path)
    fig = plt.gcf()
    fig.set_size_inches(18,10)
    plt.axis('off')
    plt.imshow(cv2.cvtColor(image, cv2.COLOR_BGR2RGB))

path = "/content/darknet/predictions.jpg"
show_detection_result(path)
```



실습

- horses.jpg에 대해 수행해 보기